

DATAFLOW 2025

MASTERING THE DATA WAVES

TEAM PROJECT FIRST ROUND

FORECASTING BUSINESS PERFORMANCE [VN]

Members:

Pham Duy Anh (leader)

Le Quoc Anh

Nguyen Thai Nhat Anh

Nguyen Lam Tung

anhpd.23BI14023@usth.edu.vn

anhlq.23BI14025@usth.edu.vn

anhntn.23BI14024@usth.edu.vn

tungnl.23BI14446@usth.edu.vn

Team No Name

2025/2026

Abstract

All of our Documents will be contained here: [Link](#)

A quick Overview of what is in our Github Repository

- The **Data** folder contains:
 - The Original Dataset stored in .xlsx file, which contains 3 sheets
 - Extracted sheets into .csv files and the Sales data already split into 'train' and 'test'
 - A .sav file that is extracted from Data Cleaning and Feature Engineering to run Statistical Analysis on SPSS
- The **Model** folder contains:
 - The Implementation and Training of ARIMA Model
 - The Implementation and Training of LSTM/RNN Model
 - The Implementation and Training of XGBoost and Tree-Based Model
- The **Preprocessing** folder contain:
 - A file that specifies how we cleaned and transformed the data into 1 final table
 - A file that compresses all of the cleaning and transforming part into 1 block and export into 1 file for model training
 - The Exported .csv file for Model training
- The **Report** folder contains:
 - The .pbix file which shows all of our Visualizations and Graphical Analysis implemented on PowerBI.

Contents

1	Introduction	4
1.1	Project Background and Scope	4
1.1.1	Background	4
1.1.2	Project Scope	4
1.2	Dataset Introduction	4
1.3	Cleaning and Transforming the Data	4
2	Data Descriptive Analysis (EDA & Visualize).	5
2.1	Overall Sale Index	5
2.2	Detailed Sale Performance	7
3	Predictive Models	9
3.1	The AutoRegression Integrated Moving Average (ARIMA) Model	9
3.1.1	Prerequisites and Problems	9
3.1.2	Determine the parameters p, d, q	11
3.1.3	Run the ARIMA and SARIMA Models with given parameter	14
3.1.4	Predict Sales for 2023	15
3.2	The Recurrent Neural Network (RNN) and Long-Short Term Memory (LSTM) Model.	17
3.3	The Extreme Gradient Boosting (XGBoost) Model	17
3.3.1	Preprocessing	17
3.3.2	Running Model and Performance Metrics	17
3.4	The Random Forest Model	17
3.4.1	Initialization and the parameter choosing steps	17
4	Results, Evaluation, and Strategy.	18
4.1	Comparison of Selected Models Performance	18
4.2	Model Evaluation and Prediction for 2023 Revenue	18
4.3	Propose upcoming investment strategies	18

List of Figures

1	Overview of Final Data	5
2	Revenue across time	5
3	Revenue of a top-performing city.	6
4	Example of revenue appearing once and vanished	6
5	Total number of cities recording revenue for less than 5 days.	6
6	Number of cities recording revenue for less than 5 days each year.	6
7	Sales map of Central Region	7
8	Top 5 City with highest Sales	7
9	Top Dates with highest sales.	7
10	Revenue and Profit Growth of each Brand	8
11	Sales index of Top 1 Brand	8
12	Revenue by Category	9
13	Revenue by Segment	9
14	New aggregated data	10
15	Data Seasonality	11
16	ACF and PACF	12
17	ARIMA model Result	15
18	Visualized Prediction	16
19	Detailed Prediction for 2023	16

1 Introduction

1.1 Project Background and Scope

1.1.1 Background

A major fashion company in the United States operates across a vast market, generating sales in over 14,000 cities across the Central, Western, and Eastern regions. As a Data Analyst in the company's Data Analysis department, your task is to analyze past index, forecast the company's future sales and provide data-driven suggestions to increase the company performance.

1.1.2 Project Scope

- Analyze sales data from the years 2010 to 2020.
- Build a forecasting model for the years 2021-2022 to assess its accuracy and make predictions for 2023.

1.2 Dataset Introduction

The original data is an .xlsx file consisting of three sheets (SalesData, Product, Geography), with the Fact table being Sales, containing 976,243 rows and 6 columns.

1.3 Cleaning and Transforming the Data

With the 3 datasets: 1 Fact and 2 Dimensions table, we can clean and merging thm into 1 main table used for training and testing models

The code of each step will be given below:

```
1 #Split into 3 different DataFrames
2 sales = df['SalesFact']
3 geo = df['Geography']
4 product = df['Product']
5
6 #Convert Date column into datetime format
7 sales['Date'] = pd.to_datetime(sales['Date'])
8
9 #Handle null in Revenue
10 sales['Revenue'] = sales['Revenue'].replace(np.nan, sales['Revenue'].median())
11
12 #Add column profit
13 sales['Profit'] = sales['Revenue'] - sales['COGS']
14
15 #Split City in Geography
16 geo['City'] = geo['City'].str.split(',').str[0]
17
18 #Split District in Geography
19 geo['District'] = geo['District'].str.split('#').str[1]
20
21 #Join Sales and Geo into 1
22 sales = pd.merge(sales, geo, on = 'Zip', how = 'left')
23
24 #Delete some columns in Sales table
25 sales = sales.drop(['State', 'District', 'Zip'], axis = 1)
26
27 #Join Sals and Product
28 sales = pd.merge(sales, product, on = 'ProductID', how = 'left')
29
30 #Get product brand
31 sales['Product_Brand'] = sales['Product'].str.split(' ').str[0]
```

```

32
33 #Delete some columns
34 sales = sales.drop(['Product', 'ProductID', 'Unnamed: 6'], axis = 1)
35
36 sales['Month'] = sales['Date'].dt.month
37 sales['Year'] = sales['Date'].dt.year
38 sales['Day'] = sales['Date'].dt.day
39
40 #Split sales into train and test
41 train = sales[sales['Date'].dt.year <= 2020]
42 test = sales[sales['Date'].dt.year >= 2021]

```

After cleaning and transforming the data, it was merged into a single table with 976,243 rows and 14 columns.

	ID	Date	Units	Revenue	COGS	Profit	City	Region	Category	Segment	Product Brand	Month	Year	Day
0	0	2013-07-31	12	19648.44	12309.747660	7338.692340	Austin	Central	Urban	Convenience	Pirum	7	2013	31
1	1	2014-03-12	16	20351.52	13497.128064	6854.391936	Torrance	West	Rural	Productivity	Natura	3	2014	12
2	2	2013-11-29	26	111367.62	91488.499830	19879.120170	Salem	East	Urban	Convenience	Currus	11	2013	29
3	3	2018-08-29	12	36280.44	21967.806420	14312.633580	Lithia Springs	East	Rural	Select	Pirum	8	2018	29
4	4	2013-04-27	14	55557.18	48645.866808	6911.313192	Troup	Central	Urban	Convenience	Natura	4	2013	27
...
976238	976238	2013-08-12	1	1070.37	749.794185	320.575815	McKeesport	East	Rural	Productivity	Aliqui	8	2013	12
976239	976239	2011-12-04	1	1070.37	749.794185	320.575815	Littleton	Central	Rural	Productivity	Aliqui	12	2011	4
976240	976240	2011-05-29	1	1070.37	749.794185	320.575815	Cumming	East	Rural	Productivity	Aliqui	5	2011	29
976241	976241	2014-04-01	1	1070.37	749.794185	320.575815	Rocheport	Central	Rural	Productivity	Aliqui	4	2014	1
976242	976242	2012-12-22	1	1070.37	749.794185	320.575815	Hurricane	East	Rural	Productivity	Aliqui	12	2012	22

976243 rows × 14 columns

Figure 1: Overview of Final Data

2 Data Descriptive Analysis (EDA & Visualize).

2.1 Overall Sale Index

Chronologically

Below is a chart illustrating the company's revenue distribution over 10 years from 2010 to 2020. It clearly shows a declining revenue trend since 2014. The company's highest revenue periods typically fall in Q2, particularly in the months of April, June, and August each year.




Figure 2: Revenue across time

In reality, after analyzing individual cities, the Revenue decline can largely be attributed to the fact that more than 4,000 cities recorded sales for only a short period (less than 5 days) before disappearing. This phenomenon occurred between 2013 and 2017, significantly contributing to the surge in total revenue during that time. Meanwhile, the revenue of major cities remained stable and did not show a noticeable decline throughout the timeline.




Figure 3: Revenue of a top-performing city.




Figure 4: Example of revenue appearing once and vanished

```
city_counts = df['City'].value_counts()

# Find the number of cities that appear only once
unique_city_count = (city_counts < 5).sum()
print(f'City appears less than 5 times: {unique_city_count} cities')

✓ 0.1s
City appears less than 5 times: 4127 cities
```

Figure 5: Total number of cities recording revenue for less than 5 days.

```
city_counts = df['City'].value_counts()

# Get cities that appear < 5 times
cities_less_than_5 = city_counts[city_counts < 5].index

# Filter the original DataFrame
filtered_df = df[df['City'].isin(cities_less_than_5)]

# Count distinct cities per year
yearly_city_count = filtered_df.groupby('Year')['City'].nunique().reset_index()

# Sort the result
yearly_city_count = yearly_city_count.sort_values(by='Year')

print(yearly_city_count)

✓ 0.2s
```

Year	City
0	2010 329
1	2011 754
2	2012 793
3	2013 833
4	2014 877
5	2015 824
6	2016 761
7	2017 444
8	2018 398
9	2019 375
10	2020 421
11	2021 379
12	2022 241

Figure 6: Number of cities recording revenue for less than 5 days each year.

- It can be observed that every year, there are cities with recorded revenue for less than 5 days. However, from 2012 to 2016, this number increased significantly—doubling compared to other years—directly explaining the unusually high revenue during this period.

Another reason that may explain the decline in revenue in recent years is the emergence of the COVID-19 pandemic, which has impacted people's lifestyles and consumption habits.

About Geographical Distribution

The states with high revenue distribution (Light → Dark, Blue → Green) are mostly located in the Southeastern part of the Central Region, with the highest concentration in California.

This is understandable due to the nature of high industrialization and quality of living in these states

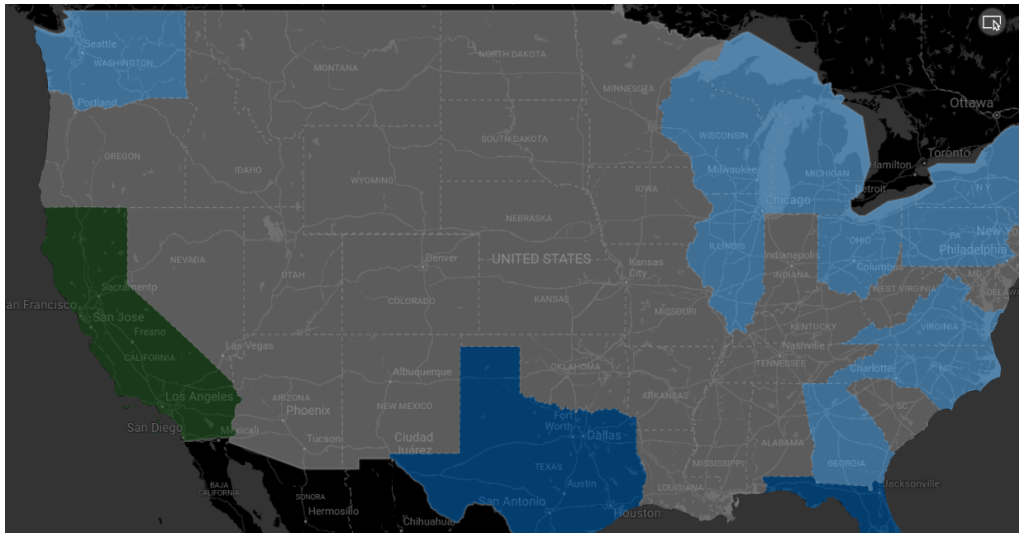


Figure 7: Sales map of Central Region

Some cities account for the company's main revenue share:

City	Revenue	Units Sold	Profit
Miami	26,472,570.39	4690	4157662.46
Houston	25,115,716.08	4190	4354357.80
Las Vegas	24,190,889.94	4268	4300447.70
San Diego	21,560,166.81	3992	3912525.64
Jacksonville	21,108,853.71	3425	3664535.67

Figure 8: Top 5 City with highest Sales

2.2 Detailed Sale Performance

June 1st (Children's Day and summer sale), August 30th (Back to school), and April 1st (Spring sale) record peak revenue each year.

Date	Units	Total Revenue
6/1/2014	3451	17,472,120.75
6/1/2016	3092	16,918,014.33
6/1/2015	3229	16,680,175.47
8/30/2014	3112	15,476,679.54
4/1/2015	2585	14,792,681.61
4/1/2013	2705	12,841,044.93
4/1/2014	2494	12,552,624.00
4/1/2012	2671	12,549,316.50
8/30/2013	2801	11,953,167.66
6/1/2013	2651	11,840,924.97

Figure 9: Top Dates with highest sales.

The Top Sales recorded date here may mostly be affected by the unexpected surge of sudden sale within cities, as listed above. The reason why they mainly focus on these 3 days will require more thorough discoveries

Sales index of Brands.

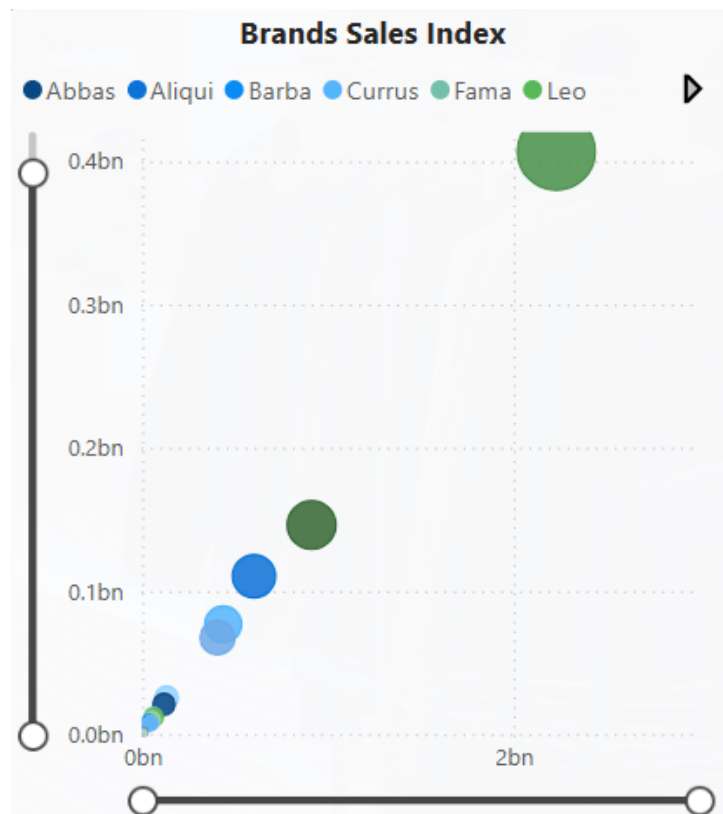


Figure 10: Revenue and Profit Growth of each Brand

The Maximus brand has been the best-selling brand over the years, surpassing all other brands. It is followed by Natura and Aliqui, whose sales are less than 50% of Maximus.

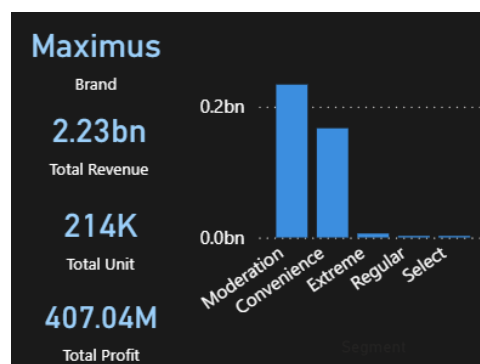


Figure 11: Sales index of Top 1 Brand

The best-selling Category and Segment over the years.




Figure 12: Revenue by Category




Figure 13: Revenue by Segment

Note: These are only the overall index of Categories and Segments, these will change depending on which Brands, timestamps or City you specifically search for

3 Predictive Models

- The presentation will use three predictive models with increasing accuracy: ARIMA → RNN/LSTM → XGBoost.
- Evaluation metrics: R-Square (R^2), Mean Absolute Percentage Error (MAPE), RMSE.

3.1 The AutoRegression Integrated Moving Average (ARIMA) Model

3.1.1 Prerequisites and Problems

Before running the ARIMA model, we need to perform clean and feature engineering 1 more time to get all relevant data to the same distinct timestamp

```

1 df_copy = df.groupby('Date').agg({
2     'Revenue': 'sum',
3     'Units': 'sum',
4     'COGS': 'sum',
5     'City': lambda x: x.mode()[0], # Most frequent value (mode)
6     'Segment': lambda x: x.mode()[0],
7     'Category': lambda x: x.mode()[0],
8 })

```

And save this as a new csv file

```

1 df_copy.to_csv('analyse_data.csv')

```


Now this step is to ensure that we bring every index into the same time structure without needing to create a rather frustrating hierarchical structure for each day, the new data will make sure that Date is distinct, Profit, COGS and Revenue are summed up and other available categorical data is taken as mode

	Date	Revenue	Units	COGS	City	Segment	Category	Year	Month	Quarter
0	2010-07-04	1765391.67	252	1.340364e+06	974	1	2	2010	7	3
1	2010-07-05	1425986.10	208	1.112580e+06	494	1	2	2010	7	3
2	2010-07-06	302463.00	33	2.373014e+05	779	1	2	2010	7	3
3	2010-07-07	1047787.65	181	8.425688e+05	727	1	2	2010	7	3
4	2010-07-08	771811.74	121	6.028101e+05	1118	1	2	2010	7	3
...
4248	2022-06-27	1625947.47	217	1.434694e+06	1080	3	2	2022	6	2
4249	2022-06-28	2042916.75	263	1.763163e+06	1015	3	2	2022	6	2
4250	2022-06-29	2284500.33	282	1.971644e+06	488	3	2	2022	6	2
4251	2022-06-30	700389.90	67	6.056699e+05	1060	3	2	2022	6	2
4252	2022-07-01	5312183.31	713	4.605905e+06	573	3	2	2022	7	3

4253 rows × 10 columns

Figure 14: New aggregated data

The new data after LabelEncoding categorical data and adding more detailed time intervals



By running the Correlation Heatmap, we can see a high multicollinearity between Revenue - Units - COGS, which is understandable, also the Time Hierarchy shows a high multicollinearity between its segments. In the final input data, only Revenue and Date will remain as Target and Timestamps, along with encoded City, Segment and Category as Exogenous features

```

1 decomposed = seasonal_decompose(main_train['Revenue'], model='additive',
2   period=365)
3 decomposed.plot()
plt.show()

```

We will do one last thing is to check for Seasonality for SARIMAX model




Figure 15: Data Seasonality

The clear, repeated fluctuations indicate a strong seasonal effect, where seasonal or recurrent monthly trends might affect how the data is distributed. Also the residuals shows unexpected features, indicating that external factors might occur

As the result, we will successively run ARIMA, SARIMA and SARIMAX model to further explore the data

3.1.2 Determine the parameters p,d,q

- Use the Augmented Dickey-Fuller (ADF) test to determine the d parameter.

```

1 def find_d(dataset):
2     i = 0
3     temp = dataset
4     while True:
5         result = adfuller(temp['Revenue'].dropna())
6         if result[1] < 0.05:
7             print(f"Number of difference: {i}")
8             print('ADF Statistic: %f' % result[0])
9             print('p-value: %f' % result[1])
10            break
11        else:
12            temp = temp.diff().dropna()
13            i = i+1
14
15            if temp["Revenue"].isna().all(): # Prevent infinite loop
16                print("All values became NaN after differencing. Check
17                    data.")
18                break

```

```

find_d(train)
✓ 0.9s

Number of difference: 1
ADF Statistic: -29.458019
p-value: 0.000000

```

It is ideal that differencing stops at 1, but that will not sufficient to clearly reduce seasonality

- Plot AutoCorrelation (ACF) and Partial AutoCorrelation (PACF) to determine the p and q parameters.

```

1 def findpq(i, dataset):
2     if i == 1:
3         X_traindiff = dataset.diff().dropna().dropna()
4     if i == 2:
5         X_traindiff = dataset.diff().dropna().diff().dropna().dropna()
6
7     fig, axes = plt.subplots(1, 2, figsize=(14, 5))
8
9     # ACF plot (for q selection)
10    plot_acf(X_traindiff, lags = 15, ax=axes[0])
11    axes[0].set_title("Autocorrelation Function (ACF)")
12
13    # PACF plot (for p selection)
14    plot_pacf(X_traindiff, lags = 15, ax=axes[1])
15    axes[1].set_title("Partial Autocorrelation Function (PACF)")
16
17    plt.show()

```

To run the ACF and PACF, we will do one more extra step to reset Date into index and use only target column 'Revenue' for the graph

```

1 train.set_index('Date', inplace = True)
2 test.set_index('Date', inplace = True)
3
4 findpq(1, train['Revenue'])

```

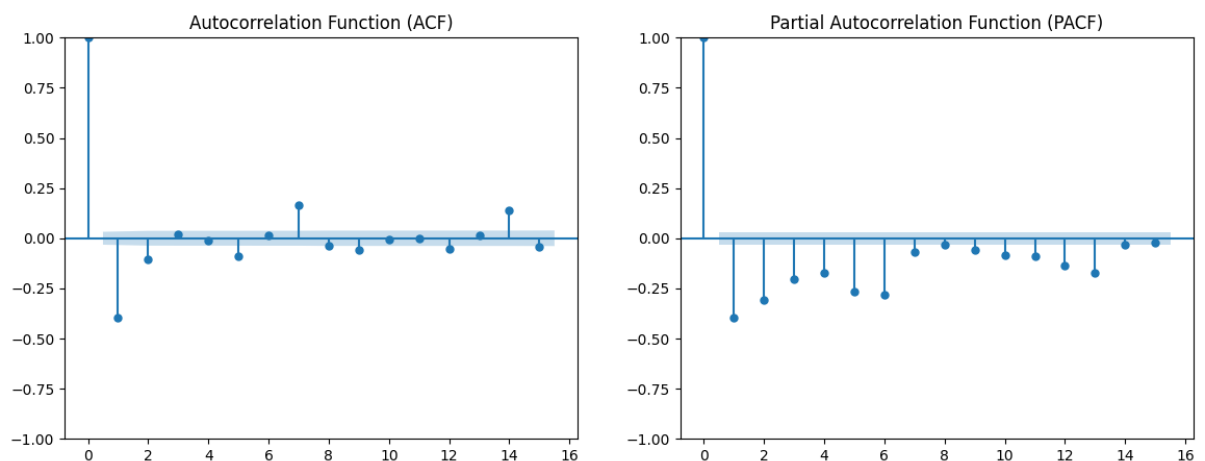


Figure 16: ACF and PACF

Take note that the Lag occur right before the sudden drop of value is the ideal lag for parameter, this case visually the optimal p and q are 1(7) and 1(7) respectively

- Use `auto_arima` to determine the optimal p, d, q parameters for the ARIMA model.

```
1 def detect(dataset):
2     auto_model = auto_arima(dataset["Revenue"], seasonal=False, trace=True)
3     print(auto_model.summary())
```

```
Best model: ARIMA(2,1,2)(0,0,0)[0]
Total fit time: 94.542 seconds

=====
SARIMAX Results
=====
Dep. Variable:          y      No. Observations:      3718
Model:                SARIMAX(2, 1, 2)  Log Likelihood      -57470.681
Date:                 Sun, 23 Feb 2025  AIC              114951.362
Time:                  11:15:54  BIC              114982.466
Sample:                0      HQIC              114962.429
                        - 3718
Covariance Type:      opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1         1.0859       0.013     84.749     0.000        1.061        1.111
ar.L2        -0.1891       0.015    -12.939     0.000       -0.218       -0.160
ma.L1        -1.9325       0.009   -212.328     0.000       -1.950       -1.915
ma.L2         0.9394       0.009    107.173     0.000         0.922         0.957
sigma2       1.709e+12   6.01e-15   2.84e+26     0.000       1.71e+12       1.71e+12
=====
Ljung-Box (L1) (Q):           0.02  Jarque-Bera (JB):       214538.82
Prob(Q):                     0.89  Prob(JB):                0.00
Heteroskedasticity (H):       0.47  Skew:                    4.43
...
```

This is the demo run of how `auto_arima` can be used to automatically track the optimal parameter, it is also surprising that the parameter p, q clearly differ from visual intuition

- Use `seasonal_compose` to check for seasonality and `auto_arima` to optimize the parameters P,D,Q,m for the SARIMA model.

```
1 def detect_parameter(train_data, p, d, q):
2     auto_arima(train_data['Revenue'], seasonal=True,
3     p=p, d=d, q=q, # We already know p,d,q
4     P=None, D=None, Q=None, # Let auto_arima find seasonal orders
5     m=7, # Time period
6     trace=True)
```

Keep in mind that the parameter $m = 7$ indicates the time interval that the model will work on, which is each week, this remains unsure whether the data will rather follow a monthly/ seasonal or yearly trend, which sets the parameter to 30, 120 and 365. That will pose a huge burden on the runtime performance and require stronger computer processor cores.

```
Best model: ARIMA(2,1,2)(2,0,2)[7]
Total fit time: 455.545 seconds
```

Tuned parameter (temporary):

$$(p, d, q, P, D, Q, m) = (2, 1, 2, 2, 0, 2, 7)$$

3.1.3 Run the ARIMA and SARIMA Models with given parameter

```

1 def arima(p,d,q, dataset, test_data):
2
3     model = ARIMA(dataset['Revenue'], order=(p,d,q))
4     model_fit = model.fit()
5     print(model_fit.summary())
6
7     start = 3031
8     pred = model_fit.predict(start = start, end= start + len(test_data) - 1,
9                             typ = 'levels')
10
11     residuals = test_data['Revenue'] - pred
12
13     rmse = root_mean_squared_error(test_data['Revenue'], pred)
14     mape = mean_absolute_percentage_error(test_data['Revenue'], pred)
15     r2 = r2_score(test_data['Revenue'], pred)
16
17     print(f"RMSE: {rmse:.2f}")
18     print(f"MAPE: {mape * 100:.2f}%")
19     print(f"R : {r2:.4f}")

```

```

1 def sarimax(p, d, q, P, D, Q, s, train_data, test_data):
2
3     # Fit SARIMAX model
4     model = SARIMAX(train_data['Revenue'],
5                     order=(p, d, q),
6                     seasonal_order=(P, D, Q, s), mle_regression = False)
7     model_fit = model.fit(dispatch=False)
8
9     # Forecast for the length of test data
10    start = 3031
11    pred = model_fit.predict(start = start, end= start + len(test_data) -
12                            1, typ = 'levels')
13
14    # Ensure prediction and test data are aligned
15    pred = pred[:len(test_data)]
16
17    print(model_fit.summary())
18
19    # Calculate evaluation metrics
20    rmse = root_mean_squared_error(test_data['Revenue'], pred)
21    mape = mean_absolute_percentage_error(test_data['Revenue'], pred)
22    r2 = r2_score(test_data['Revenue'], pred)
23
24    # Display metrics
25    print(f"RMSE: {rmse:.2f}")
26    print(f"MAPE: {mape * 100:.2f}%")
27    print(f"R : {r2:.4f}")
28
29    # Plot predictions vs actual values
30    plt.figure(figsize=(10, 5))
31    plt.plot(train_data.index, train_data['Revenue'], label='Train Data')
32    plt.plot(test_data.index, test_data['Revenue'], label='Actual Test
33             Data')
34    plt.plot(test_data.index, pred, label='Forecast', linestyle='--')
35    plt.legend()

```

```

35     plt.title('SARIMAX Forecast vs Actual')
36     plt.show()
37
38     return model_fit

```

One more time, keep in mind that the start = 3031 part in the Model code is obtained through a series of test - compare where I have to set a threshold R^2 for each step to find out the best start position. This part doesn't necessarily reflect the actual accuracy of the model but visually it does.

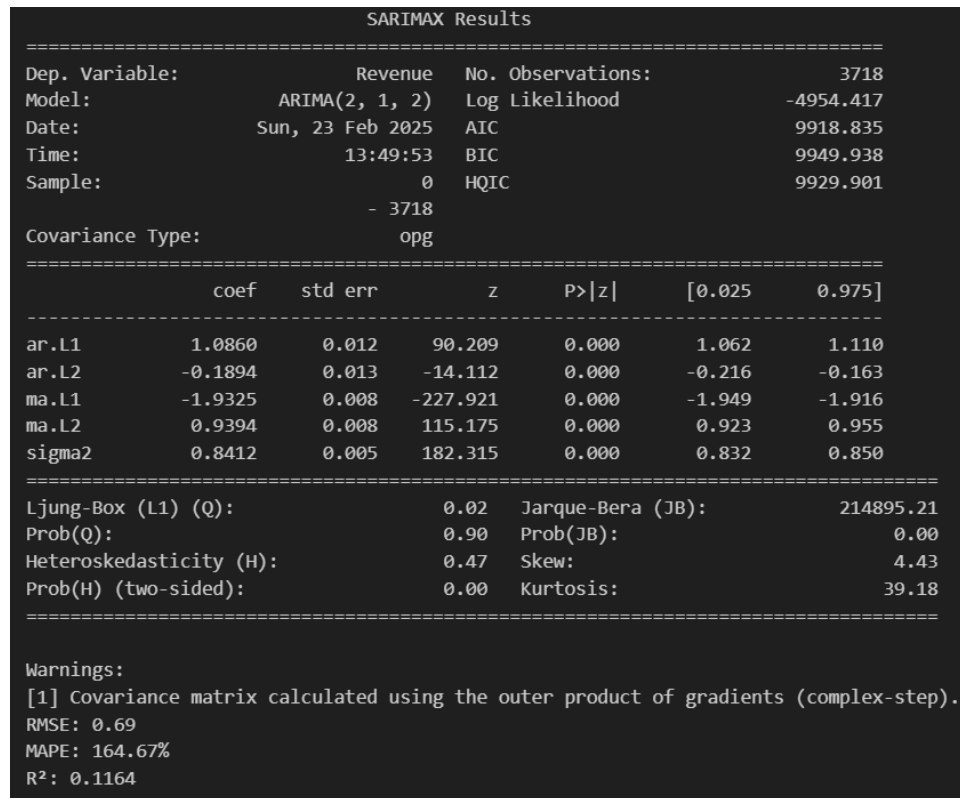


Figure 17: ARIMA model Result

3.1.4 Predict Sales for 2023

```

1 def sarimax_forecast(p, d, q, P, D, Q, s, train_data, test_data,
2     future_periods=12):
3
4     # Fit SARIMAX model
5     model = SARIMAX(train_data['Revenue'],
6                     order=(p, d, q),
7                     seasonal_order=(P, D, Q, s),
8                     mle_regression=False)
9     model_fit = model.fit(dispatch=False)
10
11     # Forecast for test data
12     start = 3031
13     pred_test = model_fit.predict(start=start, end=start + len(test_data) - 1,
14                                 typ='levels')
15
16     # Forecast for future periods
17     future_index = pd.date_range(start=test_data.index[-1],
18                                 periods=future_periods + 1, freq='M')[1:]

```



```

16 pred_future = model_fit.predict(start=start + len(test_data), end=start +
17     len(test_data) + future_periods - 1, typ='levels')
18 #pred_future = scaler.inverse_transform([pred_future])
19
20 # Calculate evaluation metrics on test data
21 rmse = root_mean_squared_error(test_data['Revenue'], pred_test)
22 mape = mean_absolute_percentage_error(test_data['Revenue'], pred_test)
23 r2 = r2_score(test_data['Revenue'], pred_test)
24
25 # Display metrics
26 print(model_fit.summary())
27 print(f"RMSE: {rmse:.2f}")
28 print(f"MAPE: {mape * 100:.2f}%")
29 print(f"R : {r2:.4f}")
30
31 # Plot predictions vs actual values
32 plt.figure(figsize=(12, 6))
33 plt.plot(train_data.index, train_data['Revenue'], label='Train Data')
34 plt.plot(test_data.index, test_data['Revenue'], label='Actual Test Data')
35 plt.plot(test_data.index, pred_test, label='Test Forecast', linestyle='--')
36 plt.plot(future_index, pred_future, label='Future Forecast (2023)',
37     linestyle='dotted', color='red')
38 plt.legend()
39 plt.title('SARIMAX Forecast vs Actual')
40 plt.show()
41
42 future = scaler.inverse_transform([pred_future])*1000
43 f = pd.DataFrame({'Month': np.arange(1,13), 'Predicted Revenue':
44     future.flatten()})
45
46 return model_fit, f

```

I've added up both SARIMA and SARIMAX into 1 final Result as the Exogenous 'City', 'Segment' and 'Category' don't bring high predictive power but rather reduce the accuracy.




Figure 18: Visualized Prediction

Month	Predicted Revenue
1	435641.557910
2	498564.981042
3	572076.255626
4	601715.841162
5	496638.456666
6	540848.192542
7	533724.080097
8	421735.676016
9	530639.008547
10	620318.290729
11	508112.582527
12	518004.285986)

Figure 19: Detailed Prediction for 2023

3.2 The Recurrent Neural Network (RNN) and Long-Short Term Memory (LSTM) Model.

3.3 The Extreme Gradient Boosting (XGBoost) Model

3.3.1 Preprocessing

3.3.2 Running Model and Performance Metrics

3.4 The Random Forest Model

3.4.1 Initialization and the parameter choosing steps

4 Results, Evaluation, and Strategy.

4.1 Comparison of Selected Models Performance

Model	Performance Metrics			Problem Involved	Future Improvement
	R2	MAPE	RMSE		
ARIMA	-0.094	112.20%	1051264	Sensitive to large datasets and has not fully optimized parameters due to high computational cost	Tune parameters further, explore data aggregation, and incorporate exogenous variables.
RNN	0.4787	21.2%	0.1236	Takes advantage of time-series but struggles to store long-term relationships. Often encounters the vanishing gradient problem, which can prevent the model from fully learning long-term trends in the data.	Adjust the number of neurons, activation functions, or add more RNN layers.
LSTM	0.5018	21.92%	0.1208	It overcomes RNNs drawback of gradually forgetting information thanks to the gating mechanism (forget, input, output). This helps the model retain long-term relationships better, leading to more favorable R^2 and RMSE metrics compared to RNN.	Apply regularization mechanisms (Dropout, L2, etc.) and adjust the architecture (multi-layer LSTM, Bi-LSTM) to enhance model generalization. Additionally, experiment with hybrid models.
XGBoost					
RandomForestRegressor					

Table 1: Summary of Model Evaluation

4.2 Model Evaluation and Prediction for 2023 Revenue

Along with the declining revenue trend in previous years, 2023 is projected to continue this downward trajectory, with estimated revenue falling within the range of:

4.3 Propose upcoming investment strategies

- Strengthen investment in Texas, Washington, Florida, and the Southeastern region (cities with historically high and stable revenue), including boosting production and launching

new products.

- Continue investing in top-performing brands such as Maximus, Natura, and Aliqui, along with best-selling products and new offerings aligned with customer preferences.
- Expand and enhance online sales platforms while developing additional marketing campaigns to increase brand awareness.
- Address key factors affecting revenue in recent years (inflation, COVID-19, emerging trends, reduced clothing expenditure) and develop strategic measures to mitigate these challenges.

References